

Performance analysis and optimization of C++ standard libraries

Aditya Kumar¹ and Sebastian Pop²

¹Samsung Austin R&D Center

²Samsung Austin R&D Center

{*aditya.k7*, *s.pop*}@samsung.com

Summary

C++ programs are widely used in performance critical applications. The standard libraries of C++, hence, are expected to be very efficient. However, experimental results show opportunities for improvements in some of the most commonly used data structures and algorithms.

We will present the performance analysis work on `libc++` and `libstdc++` and the changes we did to these libraries and to the LLVM compiler to optimize the code using them. This includes our contributions to standard library algorithms like `string::find`, `libc++::basic_streambuf::xsgetn`, and `libc++::locale`. We improved these suboptimal algorithms, particularly `string::find` which improved by more than 10x. Similarly, we enabled the inlining of constructor and destructor of `libc++::string`. We will present a systematic analysis of function attributes and the places where we added missing attributes. We will present a comparative analysis of `libc++` vs. `libstdc++` vs. Microsoft's C++ standard library on commonly used data structures and algorithms based on our `std-benchmark` (<https://github.com/hiraditya/std-benchmark>), that we started developing to help analyze standard C++ libraries. We will discuss the performance issues with `libc++::stringstream` and `libc++::sort` that we are currently working on. We will also present the lessons learned as a result of analyzing C++ standard libraries, for example:

1. Iterator based algorithms can lose information and hence, can result in suboptimal performance. This is exemplified in the implementation of `std::rotate` where we can just exchange few pointers should the underlying container is a doubly linked list e.g., `std::list`.
2. The C++ programming language has a limitation that the constructor and destructor cannot be `const` qualified which could have facilitated useful compiler optimizations like removing the destructor of a `const std::string` when the string is small enough to be kept on the stack.

Keywords: C++, performance analysis, benchmarking libraries, compiler optimization, LLVM, `libstdc++`, `libc++`

Reference to previous talks: <http://sched.co/A8J7>, <http://sched.co/8Yzk>