

# Performance analysis and optimization of C++ standard libraries

Aditya Kumar


Samsung Austin R&D Center

# Suboptimal implementation of basic\_streambuf::xsgetn

```
template <class _CharT, class _Traits>
streamsize
basic_streambuf<_CharT, _Traits>::xsgetn(char_type* __s, streamsize __n)
{
    const int_type __eof = traits_type::eof();
    int_type __c;
    streamsize __i = 0;
    for (; __i < __n; ++__i, ++__s)
    {
        if (__ninp_ < __einp_)
            *__s = *__ninp_++;
        else if ((__c = uflow()) != __eof)
            *__s = traits_type::to_char_type(__c);
        else
            break;
    }
    return __i;
}
```

# Suboptimal implementation of `basic_streambuf::xsgetn`

```
template <class _CharT, class _Traits>
streamsize basic_streambuf<_CharT, _Traits>::xsgetn(char_type* __s, streamsize __n) {
    const int_type __eof = traits_type::eof();
    int_type __c;
    streamsize __i = 0;
    while(__i < __n)
    {
        if (__ninp_ < __einp_)
        {
            const streamsize __len = _VSTD::min(__einp_ - __ninp_, __n - __i);
            traits_type::copy(__s, __ninp_, __len);
            __s += __len;
            __i += __len;
            this->gbump(__len);
        }
        else if ((__c = uflow()) != __eof)
        {
            *__s = traits_type::to_char_type(__c);
            ++__s;
            ++__i;
        }
        else
            break;
    }
    return __i;
}
```



# Performance improvements

Valgrind profile of a synthetic test case which only exercises xsgetn.

```
struct test
: public std::basic_streambuf<char> {

typedef std::basic_streambuf<char> base;
test() {}

void
setg(char* gbeg, char* gnext, char* gend) {
    base::setg(gbeg, gnext, gend);
}
};

int foo(char* input, char *output, int N) {
    test t;
    t.setg(input, input, input+N);
    char* pos = output;
    pos += t.sgetn(pos, N);
    return *pos;
}
```

	Base compiler without patch	Base compiler with patch
Total no of instructions (valgrind)	1,378,842	1,359,235
basic_streambuf::xsgetn (char*, long)	20,015	0

# Improvements to `string::find` algorithm

- Used to call the (suboptimal) generic `std::find` function
- Solution:
  - Separately implement `string::find`
  - The new algorithm gets converted to optimized versions of `memchr` and `memcmp`

# string::find original implementation

b1, e1 iterators to the haystack string

b2, e2 iterators to the needle string

```
__search(b1, e1, b2, e2) {
```

```
...
```

```
while (true)
```

```
{
```

```
    while (true)
```

```
    {
```

```
        if (__first1 == __s)
```

```
            return make_pair(__last1, __last1);
```

```
        if (__pred(*__first1, *__first2))
```

```
            break;
```

```
        ++__first1;
```

```
    }
```

```
    _RandomAccessIterator1 __m1 = __first1;
```

```
    _RandomAccessIterator2 __m2 = __first2;
```

```
    while (true)
```

```
    {
```

```
        if (++__m2 == __last2)
```

```
            return make_pair(__first1, __first1 + __len2);
```

```
        ++__m1;
```

```
        if (!__pred(*__m1, *__m2))
```

```
        {
```

```
            ++__first1;
```

```
            break;
```

```
        }
```

```
    }
```

```
}
```

```
}
```

```
...
```


```
}
```

} Find the first matching character

} Match rest of the string

# string::find new algorithm

```
inline _LIBCPP_CONSTEXPR_AFTER_CXX11 const _CharT *  
__search_substring(const _CharT *__first1, const _CharT *__last1, const _CharT *__first2, const _CharT *__last2) {  
...  
    // First element of __first2 is loop invariant.  
    _CharT __f2 = *__first2;  
    while (true) {  
        __len1 = __last1 - __first1;  
        // Check whether __first1 still has at least __len2 bytes.  
        if (__len1 < __len2)  
            return __last1;  
  
        // Find __f2 the first byte matching in __first1.  
        __first1 = _Traits::find(__first1, __len1 - __len2 + 1, __f2);  
        if (__first1 == 0)  
            return __last1;  
  
        if (_Traits::compare(__first1, __first2, __len2) == 0)  
            return __first1;  
  
        ++__first1;  
    }  
}
```



Find the first matching character

Match rest of the string

# Experimental results

Benchmark	Without patch	With patch	Gain
BM_StringFindMatch1/32768	28157 ns	2203 ns	12.8x
BM_StringFindMatch2/32768	28161 ns	2204 ns	12.8x

```
// Match somewhere towards the end
static void
BM_StringFindMatch1(benchmark::State &state)
{
    std::string s1(MAX_STRING_LEN / 2, '*');
    s1 += std::string(state.range(0), '-');

    std::string s2(state.range(0), '-');

    while (state.KeepRunning())
        benchmark::DoNotOptimize(s1.find(s2));
}
```

```
// Match somewhere from middle to the end.
static void
BM_StringFindMatch2(benchmark::State &state)
{
    std::string s1(MAX_STRING_LEN / 2, '*');
    s1 += std::string(state.range(0), '-');
    s1 += std::string(state.range(0), '*');

    std::string s2(state.range(0), '-');

    while (state.KeepRunning())
        benchmark::DoNotOptimize(s1.find(s2));
}
```



# Missing inlining opportunities in basic\_string

- Important functions not inlined.
  - `basic_string::__init(const value_type* __s, size_type __sz)`
  - `basic_string::~~basic_string()`
- Clang front end does not emit the definition of these functions (extern templates) in the IR
- Solution
  - Mark functions as inline

# Missing function attributes

- Missing `__attribute__((__noreturn__))` in important functions.
  - Prevents important compiler optimizations
  - Results in false positives in static analysis results
- `__throw.*` functions in `__locale`, `deque`, `future`, `regex`, `system_error`, `vector`

Example:

```
class __vector_base_common
{
protected:
    _LIBCPP_ALWAYS_INLINE __vector_base_common() {}
    void __throw_length_error() const _LIBCPP_NORETURN_ON_EXCEPTIONS;
    void __throw_out_of_range() const _LIBCPP_NORETURN_ON_EXCEPTIONS;
};
```

# Issues with number parsing in locale

- Uses `std::string` to store the parsed numbers
  - Results in multiple (unnecessary) calls to `memset`
- Uses suboptimal ‘find’ function to search for a character in a string (can be converted to `traits_type::find`)
- Possible characters for all kinds of numbers (octal, hex, decimal) are stored in one string
  - `__atoms = “0123456789abcdefABCDEFxX+-pPiInN”`
- Makes unnecessary copies of ‘\_\_atoms’ string which are never modified

# Issues with number parsing in locale


- Avoiding copy of `__atoms` is hard because of ABI incompatibilities.
- Current workaround is to version the change with a macro

Benchmark	Without patch	With patch	Gain
BM_lstream_numbers/32	8336 ns	7472 ns	11%

- Benchmark source:
  - `std-benchmark/cxx/stringstream.bench.cpp`
  - <https://reviews.llvm.org/D30268>


# Issues with std::sort

- Worst case  $O(N^2)$  comparisons against gcc-libstdc++  $O(N \lg N)$ 
  - PR20837




# Issues with std::sort

Time complexity (worst case)



# Issues with std::sort

Time complexity (average case)



# std-benchmark

- <https://github.com/hiraditya/std-benchmark>
  - WIP
  - Builds on Linux, Windows, Mac (Thanks to cmake)
  - Performance numbers are very stable (Thanks to google-benchmark)



# Lessons learned

- `vector::push_back` without `reserve` will cause a lot of allocations ( $\sim 2N$ )
- `vector::resize`, `string::resize` initializes the memory
  - May not be what you want
- `std::find` may not always be the right choice
  - `traits_type::find` may be very efficient for string
- Rotate but not `std::rotate` on linked lists
- The destructor of `basic_string` is difficult to optimize away

# Sequence containers

MSVC Data: 32KB	vector	list	deque
push_back	252,136	1,414,562	491,372
push_back_resize	253,664	1,338,775	402,114
push_back_reserve	252,729		

Libstdc++ Data: 32KB	vector	List	deque
push_back	60,567	403,278	47,859
push_back_resize	60,246	405,581	47,867
push_back_reserve	60,480		

Libc++ Data: 32KB	Vector	list	deque
push_back	89,629	537,019	65,522
push_back_resize	88,341	395,488	65,594
push_back_reserve	88,024		

# std::string

Data: 32KB	no_match	all_match	match1	match2	prefix
MSVC	7673	1688	9208	8998	
libc++	474	827	966	971	
libstdc++	474	849	968	971	
c string (time) /cxx string (time)					
MSVC	0.18	1.48	0.32	0.30	0.07
libc++	2.15	1332.5	73.8	69.2	0.02
libstdc++	2.14	1292.9	80.15	75.4	0.02

No\_match: no match in substring

All\_match: matches in the beginning of the string

Match1: Matches at the end

Match2: Matches at the middle

Prefix: multiple prefix matches

# Associative vs Hashed Associative (Inserting Random elements)

Data:32KB	Set	Map	Unordered_set	Unordered_map
MSVC	139	183	14	75
Libc++	102	100	21	21
Libstdc++	118	116	19	46

# compiler vs. programmer

Data:32KB	programmer	compiler	hand-optimized (memcpy)
MSVC	11,736ns	11,808ns	1,124ns
clang++	1083ns	1082ns	1478ns
g++	1084ns	1448ns	1460ns

```
const char*
assign(const char *beg,
       const char *end, char *dest) {
    while (beg != end)
        *dest++ = *beg++;
    return beg;
}
```

```
const char*
assign_res(const char * __restrict beg,
           const char * __restrict end,
           char * __restrict dest) {
    while (beg != end)
        *dest++ = *beg++;
    return beg;
}
```

# References

- <https://gcc.gnu.org/onlinedocs/libstdc++/index.html>
- [http://clang-analyzer.llvm.org/annotations.html#attr\\_noreturn](http://clang-analyzer.llvm.org/annotations.html#attr_noreturn)
- <https://reviews.llvm.org/D21103>
- <https://reviews.llvm.org/D22782>
- <https://reviews.llvm.org/D22834>
- <https://reviews.llvm.org/D21232>
- <https://reviews.llvm.org/D27068>
- <https://github.com/google/benchmark>
- <https://github.com/hiraditya/std-benchmark>

# Caution while using `std::vector`

- `push_back`:
  - Invalidates iterators
  - Causes reallocation when enough space is not available ( $\sim 2N$  space for  $N$  elements)
- inserting element(s) anywhere except the end will result in reallocation

# Alternatives to `std::vector`

- If the size is known at compile time `std::array` may be a better choice
- If reads and writes are of the same order, `std::deque` is a better choice (Find the ratio of read/write to switch the container)



# Caution with std::string

- Calls memset when resized
- The destructor of basic\_string is difficult to optimize away
- String::find does not get inlined for g++

```
#include<string>
```

```
g++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv
```

```
int main() {  
    std::string s("a");  
    s+='a';  
    return 0;  
}
```

```
clang++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv  
    call    _ZdlPv  
TODO: MSVC
```

```
#include<string>
void foo();
```

```
int main() {
    std::string s("a");
    foo()
    return 0;
}
```

```
g++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv
call    _ZdlPv
```

```
clang++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv
```

# Alternatives to std::string

```
#include<string>
```

```
int main() {  
    std::string s("a");  
    s+='a';  
    return 0;  
}
```

- **Alternative:**
  - std::array<char> when size known at compile time

# std::stable\_sort

- [https://bugs.llvm.org//show\\_bug.cgi?id=26886](https://bugs.llvm.org//show_bug.cgi?id=26886)
  - libstdc++: greater performs half as many comparisons than less for a sorted array
  - Windows: greater performs same comparisons than less for a sorted array
  - libc++: greater performs 10 times more comparisons than less for a sorted array

# Worst case time complexity vs. real world performance

time-complexity	libstdc++	libcxx
std::sort	$O(n \log n)$	$O(n^2)$
std::find	$O(n)$	$O(n)$

performance	libstdc++	libcxx
std::sort(random)	$0.58N \lg N$	$0.43N \lg N$
std::find		

# Algorithms

- `std::find` may not always be the right choice
  - `traits_type::find` may be very efficient for string
- Rotate but not `std::rotate` on linked lists

# Size of containers

Container	gcc	clang	MSVC
<code>std::vector&lt;int&gt;()</code>	24	24	24
<code>std::list&lt;int&gt;()</code>	24	24	16
<code>std::deque&lt;int&gt;()</code>	80	48	40
<code>std::set&lt;int&gt;()</code>	48	24	16
<code>std::unordered_set&lt;int&gt;()</code>	56	40	64
<code>std::map&lt;int, int&gt;()</code>	48	24	16
<code>std::unordered_map&lt;int, int&gt;()</code>	56	40	64

# Optimize for latency

Memory	Latency (cycles)
L1	4
L2	12
L3	36
RAM	36+57ns

Intel i7-4770 3.4GHz (Turbo Boost off) 22 nm. RAM: 32 GB (PC3-12800 cl11 cr2).

Source: <http://www.7-cpu.com/cpu/Haswell.html>